

Thread-Level Parallelism

CSE4100: System Programming

Youngjae Kim (PhD)

<https://sites.google.com/site/youkim/home>

Distributed Computing and Operating Systems Laboratory (DISCOS)

<https://discos.sogang.ac.kr>

Office: R911, E-mail: youkim@sogang.ac.kr

Exploiting parallel execution


- So far, we've used threads to deal with I/O delays
 - e.g., one thread per client to prevent one from delaying another
- Multi-core/Hyperthreaded CPUs offer another opportunity
 - Spread work over threads executing in parallel
 - Happens automatically, if many independent tasks
 - e.g., running many applications or serving many clients
 - Can also write code to make one big task go faster
 - by organizing it as multiple parallel sub-tasks

Today

직렬화를 방법화로 바꾸는
아라계를 작업할 수 있게 해주는
분업과 할 수 있는

- **Parallel Computing Hardware** 하드웨어
 - Multicore
 - Multiple separate processors on single chip
 - Hyperthreading
 - Efficient execution of multiple threads on single core
- **Thread-Level Parallelism** 소프트웨어
 - Splitting program into independent tasks
 - Example: Parallel summation
- **Consistency Models**
 - What happens when multiple threads are reading & writing shared state

Typical Multicore Processor



- Multiple processors operating with coherent view of memory

Benchmark Machine

- Get data about machine from `/proc/cpuinfo`
- Shark Machines
 - Intel Xeon E5520 @ 2.27 GHz
 - Nehalem, ca. 2010
 - 8 Cores
 - Each can do 2x hyperthreading

Example 1: Parallel Summation

- Sum numbers $0, \dots, n-1$
 - Should add up to $((n-1)*n)/2$
- Partition values $1, \dots, n-1$ into t ranges
 - $\lfloor n/t \rfloor$ values in each range
 - Each of t threads processes 1 range
 - For simplicity, assume n is a multiple of t
- Let's consider different ways that multiple threads might work on their assigned ranges in parallel

First attempt: psum-mutex

- Simplest approach: Threads sum into a global variable protected by a semaphore mutex.

```
void *sum_mutex(void *vargp) /* Thread routine */
{
    /* Global shared variables */
    long gsum = 0; /* Global sum */
    long nelems_per_thread; /* Number of elements to sum */
    sem_t mutex; /* Mutex to protect global sum */

    int main(int argc, char **argv)
    {
        long i, nelems, log_nelems, nthreads, myid[MAXTHREADS];
        pthread_t tid[MAXTHREADS];

        /* Get input arguments */
        nthreads = atoi(argv[1]);
        log_nelems = atoi(argv[2]);
        nelems = (1L << log_nelems);
        nelems_per_thread = nelems / nthreads;
        sem_init(&mutex, 0, 1);

        /* Create peer threads and wait for them to finish */
        for (i = 0; i < nthreads; i++) {
            myid[i] = i;
            pthread_create(&tid[i], NULL, sum_mutex, &myid[i]);
        }

        /* Check final answer */
        if (gsum != (nelems * (nelems-1))/2)
            printf("Error: result=%ld\n", gsum);

        exit(0);
    }
}
```

psum-mutex.c

psum-mutex (cont)

- Simplest approach: Threads sum into a global variable protected by a semaphore mutex.

```
/* Create peer threads and wait for them to finish */
for (i = 0; i < nthreads; i++) {
    myid[i] = i;
    pthread_create(&tid[i], NULL, sum_mutex, &myid[i]);
}

/* Check final answer */
if (gsum != (nelems * (nelems-1))/2)
    printf("Error: result=%ld\n", gsum);

exit(0);
}
```

psum-mutex.c

psum-mutex Thread Routine

- Simplest approach: Threads sum into a global variable protected by a semaphore mutex.

```
/* Thread routine for psum-mutex.c */
void *sum_mutex(void *vargp)
{
    long myid = *((long *)vargp); /* Extract thread ID */
    long start = myid * nelems_per_thread; /* Start element index */
    long end = start + nelems_per_thread; /* End element index */
    long i;

    for (i = start; i < end; i++) {
        P(&mutex);
        gsum += i;
        V(&mutex);
    }
    return NULL;
}
```

psum-mutex.c

psum-mutex Performance

- Shark machine with 8 cores, $n=2^{31}$

Threads (Cores)	1 (1)	2 (2)	4 (4)	8 (8)	16 (8)
psum-mutex (secs)	51	456	790	536	681

속도 향상 못하고 있음, 잘못된 방법

- Nasty surprise:
 - Single thread is very slow
 - Gets slower as we use more cores

Next Attempt: psum-array

- Peer thread i sums into global array element $\text{psum}[i]$
- Main waits for threads to finish, then sums elements of psum
- Eliminates need for mutex synchronization


```
/* Thread routine for psum-array.c */
void *sum_array(void *vargp)
{
    long myid = *((long *)vargp); /* Extract thread ID */
    long start = myid * nelems_per_thread; /* Start element index */
    long end = start + nelems_per_thread; /* End element index */
    long i;

    for (i = start; i < end; i++) {
        psum[myid] += i; lock free 하게 메모리, 한정된 공간을 나눠서 일부 분담해줌
    }
    return NULL;
}
```

psum-array.c

psum-array Performance

- Orders of magnitude faster than psum-mutex



Next Attempt: psum-local

- Reduce memory references by having peer thread i sum into a local variable (register)

```

/* Thread routine for psum-local.c */
void *sum_local(void *vargp)
{
    long myid = *((long *)vargp);           /* Extract thread ID */
    long start = myid * nelems_per_thread;  /* Start element index */
    long end = start + nelems_per_thread;   /* End element index */
    long i, sum = 0;


    for (i = start; i < end; i++) {
        sum += i;  /* 각각 sum 가져와 있는 것으로 보려면 → sum 을 지역변수로 한것임
                    마지막에 반환된 sum 값이옴 */
    }
    psum[myid] = sum;  /* memory access 통틀을 만다! 판권법임 한것임!
                       전에 한 방법과는 메모리를 계속 접근해서 넣어서 같음 */
    return NULL;
}

```

psum-local.c

psum-local Performance

- Significantly faster than psum-array



Characterizing Parallel Program Performance


- p processor cores, T_k is the running time using k cores
- **Def. Speedup:** $S_p = T_1 / T_p$
 - S_p is *relative speedup* if T_1 is running time of parallel version of the code running on 1 core.
 - S_p is *absolute speedup* if T_1 is running time of sequential version of code running on 1 core.
 - Absolute speedup is a much truer measure of the benefits of parallelism.
- **Def. Efficiency:** $E_p = S_p / p = T_1 / (pT_p)$
 - Reported as a percentage in the range (0, 100).
 - Measures the overhead due to parallelization

Performance of psum-local

Threads (t)	1	2	4	8	16
Cores (p)	1	2	4	8	8
Running time (T_p)	1.98	1.14	0.60	0.32	0.33
Speedup (S_p)	1	1.74	3.30	6.19	6.00
Efficiency (E_p)	100%	87%	82%	77%	75%


- Efficiencies OK, not great
- Our example is easily parallelizable
- Real codes are often much harder to parallelize

Memory Consistency



- What are the possible values printed?
 - Depends on memory consistency model
 - Abstract model of how hardware handles concurrent accesses
- Sequential consistency
 - Overall effect consistent with each individual thread
 - Otherwise, arbitrary interleaving


Sequential Consistency Example



- Impossible outputs
 - 100, 1 and 1, 100
 - Would require reaching both Ra and Rb before Wa and Wb

Non-Coherent Cache Scenario

- Write-back caches, without coordination between them




Snoopy Caches

- Tag each cache block with state

Invalid Cannot use value

Shared Readable copy


Exclusive Writeable copy



Snoopy Caches

- Tag each cache block with state

Invalid	Cannot use value
Shared	Readable copy
Exclusive	Writeable copy



`print 2`

`print 200`

- When cache sees request for one of its E-tagged blocks
 - Supply value from cache
 - Set tag to S